

Parallel Programming

Lecture 9

Prof. Dr. Jonas Posner

✉ jonas.posner@cs.hs-fulda.de

Hochschule Fulda

University of Applied Sciences



Winterterm 2025/2026
⬇ E-Learning: AI5085


12th December 2025

Organization

- **No** lecture next week 😊
 - Exercises take place as usual
- Lectures will resume in the week beginning January 12
 - I.e, Lecture 10 on January 16, 2025
 - Exercises: questions regarding the OpenMP exam

Plan for Today

- Distributed Memory
- Literature
- What is MPI?
 - Hello World with MPI
 - Compile
 - Execute




Literature—MPI

- MPI: A Message-Passing Interface Standard
 - <https://www.mpi-forum.org>
- *The same as for OpenMP:*
 - Rauber & Rünger: *Parallele Programmierung*, Springer 2012
 - Rauber & Rünger: *Parallel Programming for Multicore and Cluster Systems*, Springer, 2013
 - Hoffmann & Lienhart: *OpenMP: An Introduction to Parallel Programming with C/C++*, Springer 2008
 - Grama et al.: *Introduction to Parallel Computing*, Addison-Wesley, 2004
 - Eijkhout: *The Art of HPC*, <https://theartofhpc.com>
- *And a lot more ...*

Model: Message Passing

- Also known as **two-sided communication**
- Multiple processes execute **in parallel**
- There is **no shared memory**
 - Thus, data distribution is required.
- Communication occurs **exclusively through sending and receiving messages**
- Sending and receiving are **explicit operations** performed by both processes involved
 - (in the classical model)



Message Passing Interface (MPI) (1/3)

- Designed primarily for **multicomputer systems**
 - But efficient implementations also exist for multiprocessors
- Current version: MPI 5.0 (June, 2025)
 - MPI 1.0 in 1994
 - MPI 4.0 in 2021
- MPI is a message-passing *standard*
 - For communication between processes
 - Defines syntax and semantic of library routines
 - Portable
 - *Not an implementation!*
- Developed by academia and industry
- <https://www.mpi-forum.org>



Message Passing Interface (MPI) (2/3)

- De-facto communication standard for parallel HPC programs
- Goals: high performance, scalability, and portability
- Data exchange must be performed by explicit message-passing
- One process is executed on one CPU / Core
- Each MPI process has its own private memory and has no direct access to the memory of the other processes



Message Passing Interface (MPI) (3/3)

- MPI only partially specifies:
 - Process startup
 - Process-to-processor mapping
- Language bindings exist for **C** (C++) and **Fortran**
- Typically, the number of processes is **fixed** at program start and remains constant during execution
- Common implementations: **Open MPI**, **MPICH**, **MVAPICH**, and others



Open MPI

- *Open Source High Performance Computing*
- Implementation of the MPI standard
- Current Version: 5.0.7 (February 2025)
 - Full MPI-3.1; many elements from MPI-4.0
- Developed and maintained by a group of academic, research, and industry partners
- Language bindings: C (and Fortran)
- <https://www.open-mpi.org>




OPEN MPI

Other Models / Impls for Distributed Memory

- PGAS (Partitioned Global Address Space)
 - Parallel programming model
 - Assumes global memory, logically partitioned and a portion of it is local to each process
 - Library based: Global Arrays, OpenSHMEM
 - Compile based: Unified Parallel C (UPC), Co-Array Fortran (CAF)
 - Language-based: Chapel (Cray, HPE), X10 (IBM)

MPI: Getting Started

- Major MPI Features
 - Point to point Two sided Communication
 - Collective Communication
 - One-sided Communication
 - Parallel I/O
- Why?
 - Abstract message and file I/O exchange
 - Simplifies programming
 - Overlap computation and communication:
 - e.g. Non-blocking collectives and sending/receiving
 - Resiliency
 - Integrated failure detection



MPI—Hello World in C

```
1 #include <mpi.h>
2 #include <stdio.h>
3
4 int main(int argc, char** argv) {
5     int world_size, world_rank;
6
7     MPI_Init(&argc, &argv); // Start MPI
8     MPI_Comm_size(MPI_COMM_WORLD, &world_size); // Get total processes
9     MPI_Comm_rank(MPI_COMM_WORLD, &world_rank); // Get process ID
10
11    printf("Hello from process %d of %d\n", world_rank, world_size);
12
13    MPI_Finalize(); //Clean up MPI
14    return 0;
15 }
```

MPI—Hello World in C

```
1 #include <mpi.h>
2 #include <stdio.h>
3
4 int main(int argc, char** argv) {
5     int world_size, world_rank;
6
7     MPI_Init(&argc, &argv); // Start MPI
8     MPI_Comm_size(MPI_COMM_WORLD, &world_size); // Get total processes
9     MPI_Comm_rank(MPI_COMM_WORLD, &world_rank); // Get process ID
10
11    printf("Hello from process %d of %d\n", world_rank, world_size);
12
13    MPI_Finalize(); //Clean up MPI
14
15 }
```

- Each process runs the same code
 - *Single Program Multiple Data (SPMD)*
- Processes are identified by ranks

MPI—Hello World in C

```
1 #include <mpi.h>
2 #include <stdio.h>
3
4 int main(int argc, char** argv) {
5     int world_size, world_rank;
6
7     MPI_Init(&argc, &argv); // Start MPI
8     MPI_Comm_size(MPI_COMM_WORLD, &world_size); // Get total processes
9     MPI_Comm_rank(MPI_COMM_WORLD, &world_rank); // Get process ID
10
11    printf("Hello from process %d of %d\n", world_rank, world_size);
12
13    MPI_Finalize(); //Clean up MPI
14    return 0;
15 }
```

Compile: mpicc -o hello hello.c
Run: mpirun -np 4 ./hello

MPI—Hello World in C

```
1 #include <mpi.h>
2 #include <stdio.h>
3
4 int main(int argc, char** argv) {
5     int world_size, world_rank;
6
7     MPI_Init(&argc, &argv); // Start MPI
8     MPI_Comm_size(MPI_COMM_WORLD, &world_size);
9     MPI_Comm_rank(MPI_COMM_WORLD, &world_rank);
10
11    printf("Hello from process %d\n", world_rank);
12
13    MPI_Finalize(); //Clean up MPI
14
15 }
```

Compile: mpicc -o hello hello.c
Run: mpirun -np 4 ./hello

Output is in non-deterministic order:
Hello from process 2 of 4
Hello from process 0 of 4
Hello from process 3 of 4
Hello from process 1 of 4

What Happens During Execution?

- Single Program, Multiple Data (SPMD)
- The program is started by an **MPI launcher**
 - E.g. `mpirun` or `srun` (old: `mpiexec`)
 - The launcher starts **multiple identical processes**
- Each process:
 - Calls `MPI_Init()` to join the global communicator
 - Receives a unique **rank ID** from 0 to $N-1$
 - Can query the total number of processes
 - Prints its own message independently
- Output order is **non-deterministic**
 - Processes run concurrently
- *No sending and receiving of messages*

MPI—Compilation and Execution

- Experimental environments:
 - Fulda HPC Cluster, LinuxLab machines, or your local setup
- **Compilation:**
 - `mpicc -O3 -Wall -o hello_mpi hello_mpi.c`
 - (for C++: use `mpiCC` or `mpic++`)
- **Interactive execution:**
 - `mpirun -n 4 ./hello_mpi`
 - Launches 4 processes that run in parallel and communicate via MPI
 - *Live Demonstration*

Basic MPI Functions

- Core MPI routines:
 - MPI_Init
 - MPI_Finalize
 - MPI_Comm_size
 - MPI_Comm_rank
 - MPI_Send
 - MPI_Recv
- All library functions, datatypes, and constants start with the prefix MPI_ and are declared in mpi.h
- On success, functions return MPI_SUCCESS
 - Otherwise, the program will typically **abort by default**

MPI_Init and MPI_Finalize


- `int MPI_Init(int *argc, char ***argv)`
 - Initializes the MPI runtime environment
 - **No** MPI functions may be called **before** `MPI_Init()`
 - The arguments correspond to the program's command-line parameters
 - https://docs.open-mpi.org/en/v5.0.x/man-openmpi/man3/MPI_Init.3.html
- `int MPI_Finalize()`
 - Deregisters the calling process from the MPI runtime system
 - Does **not** terminate the program or process itself
 - After this call, **no further MPI functions** may be used
 - May block until all outstanding MPI operations are completed
 - https://docs.open-mpi.org/en/v5.0.x/man-openmpi/man3/MPI_Finalize.3.html

Communicators

```
1 int MPI_Comm_rank(MPI_Comm comm, int *rank);  
2 int MPI_Comm_size(MPI_Comm comm, int *size);
```

- A **communicator** (`comm`) is an internal, distributed MPI data structure that manages a set of processes
 - Each process holds a local part of this data structure, containing information about how to reach all other processes
- Within each communicator:
 - Processes are numbered starting from 0.
 - `size` = total number of processes in the communicator.
 - `rank` = the process's own ID
- Currently, we use only: `comm = MPI_COMM_WORLD`
 - The communicator including **all processes**

Point-to-Point-Communication



Has to know:

- **Whom to send the data to?**
Receiver's rank
- **What kind of data to send?**
E.g., 100 integers
- **Message tag:**
User-defined label (like an email subject)

Might have to know:

- **Who is sending?**
MPI_ANY_SOURCE allowed
- **What is received?**
E.g., up to 1000 int
- **Tag of the message?**
MPI_ANY_TAG allowed

Sending Data (MPI_Send)

```
1 int MPI_Send(void *buf, int count, MPI_Datatype datatype,  
2               int dest, int tag, MPI_Comm comm);
```

- buf = start address of the send buffer
 - Must be declared, allocated, and filled beforehand
- count = number of data elements to send
 - Measured in units of datatype)
- datatype = type of each message element
 - E.g. MPI_INT, MPI_FLOAT, MPI_CHAR, MPI_UNSIGNED, MPI_BYTE
- dest = rank of the **destination process**
 - Must always be specified
- tag = identifier or label
 - Freely chosen by the programmer for application-specific use

Receiving Data (MPI_Recv) (1/2)

```
1 int MPI_Recv(void *buf, int count, MPI_Datatype datatype, int source,  
2                 int tag, MPI_Comm comm, MPI_Status *status);
```

- buf = start address of the receive buffer
 - Must be declared and allocated before the call
- count = upper bound on the number of elements to receive
- datatype = datatype of the expected message elements
- source = rank of the sender, or MPI_ANY_SOURCE
- tag = specific tag value, or MPI_ANY_TAG

Receiving Data (MPI_Recv) (2/2)

```
1 int MPI_Recv(void *buf, int count, MPI_Datatype datatype, int source,  
2                 int tag, MPI_Comm comm, MPI_Status *status);
```

- status = information about the received message:
 - status.MPI_SOURCE and status.MPI_TAG
- To retrieve the actual message length:
 - int MPI_Get_count(MPI_Status *status, MPI_Datatype datatype, int *count);
- Alternatively: use MPI_Recv(..., MPI_STATUS_IGNORE) if message details are not needed

Matching Sending and Receiving

rank s calls:

```
MPI Send(send_buf, send_buf_size, send_type, dest, send_tag, send_comm);
```

rank q calls:

```
MPI Recv(recv_buf, recv_buf_size, recv_type, src, recv_tag, recv_comm, &status);
```

- All 5 **green** parameters need to match to get message successfully through
 - All mandatory to be equal,
except `recv_buf_size >= send_buf_size`


MPI: Elementary Datatypes

MPI_datatype	C equivalent
MPI_SHORT	short int
MPI_INT	int
MPI_LONG	long int
MPI_LONG_LONG	long long int
MPI_UNSIGNED_CHAR	unsigned char
MPI_UNSIGNED_SHORT	unsigned short
MPI_UNSIGNED	unsigned int
MPI_UNSIGNED_LONG	unsigned long
MPI_UNSIGNED_LONG_LONG	unsigned long long
MPI_FLOAT	float
MPI_DOUBLE	double
MPI_LONG_DOUBLE	long double
MPI_BYTE	char

- C types cannot be passed directly
 - Use **MPI datatypes** instead
- Advantage:
 - Interoperability with heterogeneous software and hardware

Conditions for Message Exchange

- Communicators must be **identical**, and all other parameters must be **compatible**
- Messages from the same sender to the same receiver are **not reordered**
- However, consider the following scenario:
 - P1: Send m_1 , then Send m_2
 - P3: Recv m_2 , then Send m_3
 - P2: Recv from ANY_SOURCE
→ **Either m_1 or m_3 may be received!**
- If matching Send/Recv operations are posted, MPI guarantees that **at least one** of them will complete (progress)
- **Fairness is not guaranteed!**



Standard Mode Communication (1/2)

- MPI_Send and MPI_Recv perform **blocking communication**:
 - After MPI_Send returns, the send buffer may be safely reused
 - After MPI_Recv returns, the receive buffer contains the rec. data
- In **standard mode**, returning from MPI_Send may indicate:
 - (1) The message has been delivered into the receiver's buffer, **or**
 - (2) The message has been **temporarily buffered** by the system
- Whether (1) or (2) occurs depends on:
 - The MPI implementation
 - The current system state (e.g., whether the receive is already posted, or if buffer space is available)

Standard Mode Communication (2/2)

- A program is called **safe** if its correctness does not depend on the existence or size of MPI system buffers
 - In particular, **no deadlocks** must occur regardless of buffering

Again: Hello World

```
1 MPI_Init(&argc, &argv);
2 int rank, size;
3 MPI_Comm_rank(MPI_COMM_WORLD, &rank);
4 MPI_Comm_size(MPI_COMM_WORLD, &size);
5 int mes;
6 if (rank == 0) {
7     mes = 42; // start value
8     MPI_Send(&mes, 1, MPI_INT, 1, 0, MPI_COMM_WORLD);
9 } else if (rank == size - 1) { // last rank does not forward
10     MPI_Recv(&mes, 1, MPI_INT, rank - 1, 0, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
11 } else {
12     MPI_Recv(&mes, 1, MPI_INT, rank - 1, 0, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
13     mes++; // modify payload
14     MPI_Send(&mes, 1, MPI_INT, rank + 1, 0, MPI_COMM_WORLD);
15 }
16 printf("Hello from rank %d/%d, message = %d\n", rank, size, mes);
17 MPI_Finalize();
18 return 0;
```

Rules for executing MPI programs on the cluster

- **Slurm should be used for any program execution!!!**
 - **No** local executions on the login node!
- Execute OpenMP Programs always on exact one node (64 Hardware-Threads)
- Execute MPI Programs with a maximum of 832 Processes
 - The cluster comprises 13 nodes (node-00—node-12), each with 64 physical cores (no hyperthreading!)
 - ⇒ Maximum number of cores: 832
 - ⇒ Start your MPI programs with a maximum of 832 processes
 - If you are using EduMPI, please do not exceed the maximum of 400 processes