

Parallel Programming

Lecture 12

Prof. Dr. Jonas Posner

✉️ jonas.posner@cs.hs-fulda.de

Hochschule Fulda

University of Applied Sciences




Winterterm 2025/2026
⬇️ E-Learning: AI5085

30th January 2026


Plan for Today

- **Register for a presentation slot!**
 - Deadline: January 30 at 11:59 p.m.
- Introduction to Performance Analysis
 - Concepts and terminology
 - Tools and workflow
 - Get into practice

Experiment: Ping Pong—Blocking vs. Non-Blocking



VS.



- `PING_PONG_LIMIT = 100`
- Computation = 10 ms
- Two processes

Variant	Overlap?	Measured Time	Speedup
Blocking	✗ No	≈ 1.00 s	1
Non-Blocking	✓ Yes	≈ 0.50 s	2

- How to “see” this in real programs? (From Lecture 10)
Today!

Motivation for Performance Analysis

- Goals for high-quality parallel programs:
 - **Correctness:** no races, deadlocks, or mismatched communication
 - **Efficiency:** high utilization of available resources
 - **Scalability:** efficiency holds as problem size and process count grow
- Performance analysis tools help to:
 - Identify bottlenecks and inefficiencies
 - Validate optimizations with empirical data
 - Understand complex parallel behavior

Instrumentation

- **Instrumentation:** Add code to collect performance data
 - Manually (developer)
 - Automatically (compiler)
 - Via pre-instrumented libraries (e.g., MPI wrappers)
- **Static instrumentation:** Modify code at compile time
 - Analysis after execution (*post-mortem*)
 - Recompile to change instrumentation
- **Dynamic instrumentation:** Modify at runtime
 - Operates directly on the binary
 - No recompilation; flexible but potentially higher overhead

Profiling

- Profiling collects **summary statistics** about program behavior
 - **Event-based:** Metrics per event (e.g., function calls)
 - **Sampling-based:** Periodic interrupts capture program state
- Typical metrics in MPI programs:
 - Number and duration of MPI calls
 - Time in computation vs. communication
- Profile data is **aggregated**
 - Temporal ordering is not preserved
 - Good overview, not a detailed timeline
- Lower overhead than tracing; suitable for long-running applications

Tracing

- Tracing records **timestamped events** during execution
 - Captures *when* and *where* each event happened
 - Preserves temporal ordering and causality
- Trade-offs compared to profiling:
 - Higher storage requirements and runtime overhead
 - More detail; supports root cause analysis
 - Requires specialized tooling for analysis

Example

Trace:

```
[2021-06-12T11:22:09.815479Z] [INFO] [Thread-1] Request started
[2021-06-12T11:22:09.935612Z] [INFO] [Thread-1] Request finished
[2021-06-12T11:22:59.344566Z] [INFO] [Thread-1] Request started
[2021-06-12T11:22:59.425697Z] [INFO] [Thread-1] Request finished
```

Profile:

```
2 "Request finished" events
2 "Request started" events
```

Performance Analysis

- Based on data collected during instrumentation and measurement
- Targets for identification:
 - Inefficient code sections
 - Communication bottlenecks
 - Load imbalances
 - Resource contention (e.g., oversubscribed cores)
 - Suboptimal parallelization patterns
- Analysis is often performed *post-mortem*
 - Enables detailed examination with minimal runtime interference

Feedback and Visualization

- Presents analysis results to the developer
- Common feedback mechanisms:
 - Source code annotations highlighting problematic regions
 - Automatically generated reports (text, HTML, PDF)
 - Interactive graphical user interfaces (GUIs)
- Effective visualization enables:
 - Rapid identification of performance issues
 - Correlation between metrics and source code
 - Comparison across different runs or configurations

Debugging Parallel Programs

- Goal: identify and diagnose program errors
- Common errors in parallel programs:
 - **Race conditions:** Non-determinism due to unsynchronized access
 - **Deadlocks:** Circular wait dependencies prevent progress
- MPI-specific issues:
 - **Mismatched send/receive:** message size, datatype, tag, communicator
 - Buffer-related problems (e.g., resource exhaustion)


Debugging Tools

- Commercial debuggers for parallel MPI applications:
 - ARM DDT (part of ARM Forge)
 - TotalView (Perforce)
- Both offer comprehensive feature sets:
 - Multi-process and multi-threaded debugging
 - GPU debugging support (CUDA, OpenCL)
 - Memory debugging and leak detection
- Open-source alternatives (with limited features):
 - MUST (MPI correctness checking)
 - STAT (stack trace analysis)
 - ISP (MPI verification)

Instrumentation and Measurement Tools

- Focus on data collection; separate tools needed for analysis and visualization
- **Dyninst**
 - Widely-used API for dynamic binary instrumentation
 - Foundation for tools like TAU and OpenSpeedShop
 - Enables runtime analysis without recompilation
- **Score-P**
 - Popular tool for static instrumentation
 - Unified measurement infrastructure for multiple analysis tools
 - *Details on following slide*

Instrumentation and Measurement Tools



Score-P: <https://scorepci.pages.jsc.fz-juelich.de/scorep-pipelines/docs/scorep-6.0/html/index.html>

Score-P: Scalable Performance Measurement

- *Scalable Performance Measurement Infrastructure for Parallel Codes*
- Output formats:
 - Call-path profiles (CUBE4 format)
 - Event traces (OTF2 format)
 - Compatible with Cube, Vampir, Scalasca
- Open Source; portable to “all” HPC systems
- Maintained by Jülich Supercomputing Centre (JSC)
- Usage:
 - Compile: `scorep mpicc -O2 pingpong.c -o pingpong`
 - Run: `mpirun -np 2 ./pingpong`
- <https://www.score-p.org>



Cube and Vampir: Visualization Tools

- **Cube**

- Interactive visualization of CUBE4 profiles
- Three data views (Metric, Program, and System)
- Open Source; maintained by JSC
- <https://www.scalasca.org/software/cube-4.x>



- **Vampir**

- Interactive visualization of OTF2 traces
- Timeline view of application activities and communication
- Commercial
- <https://www.vampir.eu>





Scalasca: Automated Trace Analysis

- Scalable trace-based performance analysis toolset
- **Key feature:** Automatic pattern detection
 - Identifies common inefficiency patterns (e.g., Late Sender/Receiver)
 - Quantifies wait-state overhead
- Open Source; maintained by JSC
- Workflow (requires Score-P instrumentation):
 - Run: scalasca -analyze -t mpirun -np 2 ./pingpong
 - Analyze: scalasca -examine ./scorep_pingpong_trace
- <https://www.scalasca.org/scalasca/software/scalasca-2.x>



Vampir and Scalasca




Scalasca:


<https://www.scalasca.org/scalasca>

Vampir: <https://vampir.eu>

Scalasca: Ping Pong (Blocking)—Late Sender



Scalasca: Ping Pong (Non-Block)—No Late Sender





Late Sender and Late Receiver Problems

- Communication-induced idle time due to timing mismatches
- **Late Sender**
 - Receiver blocks waiting for a message not yet sent
 - Common with blocking MPI_Recv
 - Causes receiver-side idle time
- **Late Receiver**
 - Sender may block if the receiver has not posted a matching receive
 - Typical with synchronous sends (MPI_Ssend) or rendezvous protocol (large messages)
 - Can lead to buffer pressure/exhaustion with buffered sends (MPI_Bsend) or heavy eager buffering
 - Both patterns are detectable via trace analysis (e.g., Scalasca)

All-In-One Tools

- Combine instrumentation, measurement, analysis, and visualization
- **TAU (Tuning and Analysis Utilities)**
 - Multiple instrumentation options (source, compiler, dynamic)
 - Supports MPI, OpenMP, CUDA, OpenCL, and hybrid models
 - Post-mortem visualization in 2D and 3D
 - Topology view: maps performance data to system architecture
 - Open Source; developed at University of Oregon



TAU: <http://www.cs.uoregon.edu/research/tau/home.php>

Summary

- Communication overhead significantly impacts parallel performance
 - Blocking communication introduces synchronization-induced idle time
 - Non-blocking communication enables computation-communication overlap
- Tools like Score-P, Cube, Vampir, and Scalasca help identify performance bottlenecks
 - Visualization of timings, communication patterns, etc.
 - Identification of Late Sender / Late Receiver problems

Summary

- Communication overhead significantly impacts parallel performance
 - Blocking communication introduces synchronization-induced idle time
 - Non-blocking communication enables computation-communication overlap
- Tools like Score-P, Cube, Vampir, and Scalasca help identify performance bottlenecks
 - Visualization of timings, communication patterns, etc.
 - Identification of Late Sender / Late Receiver problems

Performance optimization is an iterative process:

→ Measure, understand, optimize, repeat